

STRATEGIC USAGE

A Thesis

Presented to the Faculty of the Graduate School

of Cornell University

in Partial Fulfillment of the Requirements for the Degree of

Master of Science

by

Eliot Seo Shekhtman

May 2024

© 2024 Eliot Seo Shekhtman
ALL RIGHTS RESERVED

ABSTRACT

Real-world systems often involve some pool of users choosing between a set of services. With the increase in popularity of online learning algorithms, these services can now self-optimize, leveraging data collected on users to maximize some reward such as service quality. On the flipside, users may strategically choose which services to use in order to pursue their own reward functions, in the process wielding power over which services can see and use their data. Extensive prior research has been conducted on the effects of strategic users in single-service settings, with strategic behavior manifesting in the manipulation of observable features to achieve a desired classification; however, this can often be costly or unattainable for users and fails to capture the full behavior of multi-service dynamic systems. As such, we analyze a setting in which strategic users choose among several available services in order to pursue positive classifications, while services seek to minimize loss functions on their observations. We focus our analysis on realizable settings, and show that naive retraining can still lead to oscillation even if all users are observed at different times; however, if this retraining uses memory of past observations, convergent behavior can be guaranteed for certain loss function classes. We provide results obtained from synthetic and real-world data to empirically validate our theoretical findings.

BIOGRAPHICAL SKETCH

Eliot Seo Shekhtman is a master of science student in computer science at Cornell University, working with Professor Sarah Dean on learning theory for dynamic systems. He has been conducting machine learning research for the past four years resulting in three publications and a workshop paper—his trajectory taking him from applied computational biology to theoretical research. Interspersed in his studies, he has additionally worked two internships with Goldman Sachs and participated in part-time contract work with them, building NLP models and pipelines.

His first publication was on an online implementation of a Gaussian copula model and its applicability to the task of data imputation with Professor Madeleine Udell and it was accepted to AAAI 2022, while his second publication concerned the application of continuous diffusion to the text generation task with Professor Kilian Weinberger; it was accepted to NeurIPS 2023. His graduate studies presented him an opportunity to take coursework in machine learning theory and feedback systems which he had neglected before. This led him to join the lab of Sarah Dean, where he completed this project characterizing the convergence of multiple learners in the presence of adversarial users, where users may pursue a positive classification by strategically choosing which learners to use. This project has just been accepted to AISTATS 2024 and a shortened version of this work has been accepted and was presented at the AAAI 2024 recommendation ecosystems workshop (EcoSys 2024).

This document is dedicated to my friends, my family, and my advisors over the years.

ACKNOWLEDGEMENTS

I would like to thank Professor Igor Lednev, Professor Marlen Z. Gonzalez, Professor Madeline Udell, Professor John Whitman, and Professor Sarah Dean. In industry, I would like to thank my managers Dr. Debarun Kar and Dr. Howard Karloff. They were all extremely influential in my career path, both teaching me and letting me explore as I found a niche that felt right for me.

This work was supported by NSF CCF 2312774, NSF OAC-2311521, a LinkedIn Research Award, and a gift from Wayfair.

TABLE OF CONTENTS

Biographical Sketch	iii
Dedication	iv
Acknowledgements	v
Table of Contents	vi
List of Figures	vii
1 Introduction	1
1.1 Strategic Classification	1
1.2 Endogeneous Distribution Shift	2
1.3 Usage Choices	2
2 Problem Setting	4
2.1 Strategic Users	5
2.2 Learning Updates for Services	8
2.3 Interaction Dynamics	10
3 Theoretical Results	13
3.1 Importance of Memory	14
3.2 Characterizing Invariance	15
3.3 Convergence	16
4 Experiments	18
4.1 EXPERIMENTS	18
4.2 Five Points Dataset	18
4.3 Banknote Authentication	20
4.4 Bank Account Fraud	24
5 Relaxing Assumptions	28
5.1 Round Robin Updates	28
5.2 Non-Realizability	30
5.3 Sampled Users	32
6 Conclusion and Discussion	34
A Missing Proofs	35
A.1 Oscillations and Counter-Examples	35
A.2 Invariance	37
A.3 Convergence	42
A.4 Round Robin Updates	44
A.5 Non-Realizability	45

LIST OF FIGURES

3.1	Five datapoint example described in the proof of Proposition 1. Negative points are represented by $-$ and positive points by $+$, where bold-face indicates two overlapping points. The dashed and solid lines represent the oscillating classifiers, with the dotted line representing a zero-loss classifier.	14
4.1	5-Points dataset; the top three graphs give the $p = 0$ case while the bottom three give $p = 0.5$. Service loss is calculated after the user update but before the service update, and usages are displayed for each of the five points with the middle graphs giving the usages for model $j = 0$ and the right graphs giving the usages for model $j = 1$	19
4.2	5-Points dataset; the top three graphs give the $p = 0.1$ case while the bottom three give $p = 1.0$	20
4.3	Banknote Authentication dataset; each graph gives the positive and negative usages of each of the five models; triangle markers above the lines indicate positive usage while below the lines indicate negative, with colors giving which model the line refers to. The left graph gives the no-memory $p = 0$ setting, while the graph on the right gives the $p > 0$ setting. Model order, and hence their colors, are meaningless due to the random initialization.	22
4.4	Banknote Authentication dataset; ablating on the seed (s). The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.	23
4.5	Banknote Authentication dataset; ablating on m . The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.	24
4.6	Banknote Authentication dataset; ablating on usage cost power factor q . The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.	25
4.7	Bank Account Fraud dataset; each graph gives the positive and negative usages of each of the five models. Notation is the same as that for Figure 4.3.	26
4.8	Bank Account Fraud dataset; ablating on the seed (s). The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.	27

5.1	Bank Account Fraud dataset in the unrealizable setting. The left plot gives the $p = 0$ case, the middle plot gives the $p = 0.5$ case using Equation 2.4 for the memory update, and the right plot gives the $p = 0.5$ case using Equation 5.1 for the memory update.	31
5.2	Banknote Authentication dataset in the sampled users setting.	33

CHAPTER 1

INTRODUCTION

1.1 Strategic Classification

Our work is inspired by the setting of *strategic classification*, first proposed by (13), in which strategic users manipulate their features in order to receive a positive classification, and a learner or decision maker is tasked with designing a classifier robust to these manipulations. Many works in this setting study complexities such as the distinction between gaming and improvement (19), connections to causal inference (22), and the social burden (23). As this form of gaming is distinct from the one we study, we do not attempt a comprehensive review of this vast body of work; instead, we highlight work that considers phenomena relevant to our usage setting. One such phenomenon is decision-dependent access to data, which in our setting arises due to user choices between services. (14) studies an online variant of the strategic classification problem in the presence of “apple tasting” or one-sided feedback, in which labels are only collected for data points that are positively classified. (5) refer to this phenomenon as algorithmic censoring and explore its implications. In contrast to these works, in the strategic usage setting, both features and labels are unavailable to services that a user does not select. Another phenomenon of interest is repeated interaction between learning algorithms and strategic agents. (7) present algorithms for an online variant of strategic classification in which data arrives sequentially, each point responding to the currently deployed classifier. (32) study the dynamics of repeated interactions between a decision-maker and a strategic population, and show convergence to a unique equilibrium depending on their relative update frequencies. Interestingly, they show that repeated retraining is sufficient to counteract manipulations when their update frequencies are high enough. Though we

study similar repeated retraining dynamics, our analysis differs in that equilibria are not unique.

1.2 Endogeneous Distribution Shift

The dynamics of repeated interactions between learning algorithms and endogenously shifting populations have also been studied more generally. *Performative prediction*, first introduced by (27), generalizes the setting of strategic classification, modeling a single learner seeking to maximize accuracy subject to an underlying decision-dependent data distribution. They also study the convergence of repeated retraining. (25; 28; 29) study a multi-player scenario, where the data distribution depends on the decisions of multiple learners. We also consider decision-dependent distributions to model the varying usage of strategic individuals; however, the mechanics of the dependence violate assumptions necessary to apply previous work in the performative setting, such as distribution smoothness. In contrast the unique equilibrium of performative prediction dynamics, our setting requires a careful convergence analysis in terms of invariant sets rather than single points. The framework of *performative power* introduced by (author?) (12) studies the ability of services to influence data distributions. Interestingly, they show that in the presence of a choice between competing providers, individuals have no incentive to perform costly manipulations to their features.

1.3 Usage Choices

A largely separate body of work has investigated the impacts of ML by studying user participation choices. (15; 31) consider sub-populations choosing whether or not to use

a single ML model on the basis of accuracy or performance, showing that retention dynamics can lead to the exacerbation of disparate representation found in the population. (10; 20; 6) consider users selecting among various services, also on the basis of model accuracy. For such non-strategic usage, these works characterize the convergence of a simple repeated retraining dynamic. We show that when usage decisions are made strategically, naive repeated retraining may fail to converge. Another line of work considers explicit competition for market share between multiple providers, where the market share is determined by users selecting based on performance or accuracy (11; 1; 17; 2; 3). While these works investigate strategic behavior on the part of services, our focus is on strategic behaviors by users. Our setting departs from all works mentioned in this subsection in that we model users who seek positive classifications, rather than merely high accuracy.

CHAPTER 2

PROBLEM SETTING

We study the interactions between $n \in \mathbb{N}_+$ individuals, which we refer to as *users*, and $m \in \mathbb{N}_+$ machine learning-based services. Each user $i \in \{1, \dots, n\}$ is represented by features $x_i \in \mathcal{X}$, which encode the information about user i that is available at decision time. For example, in a digital credit example, this information may include data about an applicant’s mobile phone usage, financial transactions, location information, and social media use, among other details (8). Also associated with each user i is a label $y_i \in \{+1, -1\}$ indicating an outcome of interest. We refer to $+1$ as a positive label, which is generally seen as desirable, and -1 as a negative label, which is generally seen as undesirable. Unlike the features, the label is not visible at decision time. In the simplified digital credit example, the label corresponds to whether an individual has the financial resources to repay a loan on time; however, users need not correspond directly to human individuals. For example, in simulation experiments, we use the Banknote Authentication dataset (21). In this setting, a user corresponds to a banknote, the features are defined by a processed image of the banknote, and the label is whether or not the banknote is authentic.

Users interact with services. Each service $j \in \{1, \dots, m\}$ selects a classifier $h_j : \mathcal{X} \rightarrow \{+1, -1\}$ from some model class \mathcal{H} . This classifier predicts the unseen label of a user, given their features. As with labels, a positive prediction of $+1$ is seen as desirable, while a negative prediction of -1 is seen as undesirable. In the digital credit example, a prediction determines whether a credit service offers a loan to an individual. In the banknote example, a prediction determines whether the financial transaction is accepted.

We study the *realizable* setting, meaning that we assume there exists a classifier $h \in \mathcal{H}$ that perfectly classifies all users simultaneously, $y_i = h(x_i)$ for $i = 1, \dots, m$. This

assumption, formally stated in Assumption 3, is consistent with modern machine learning practice: expressive high-parameter model classes and high-dimensional data allow for state-of-the-art dataset interpolation (30). Several works in strategic classification also use this setting (24; 9).

Example 1. For a given feature transformation $\varphi : \mathcal{X} \rightarrow \mathbb{R}^d$, the linear model class is defined as

$$\mathcal{H} = \left\{ h(x) = \begin{cases} +1 & \theta^\top \varphi(x) > 0 \\ -1 & \theta^\top \varphi(x) \leq 0 \end{cases} \quad s.t. \quad \theta \in \mathbb{R}^d \right\}.$$

Unlike in a classical supervised machine learning setting, or even in the usual strategic classification setting (13; 32), we do not assume that services have immediate access to data about the users. Instead, the data observed by services depends on the strategic choices of the users. The following sections describe the user behavior and the learning updates of the services.

2.1 Strategic Users

Strategic users seek positive classifications, as these correspond to desirable outcomes. This desire is independent of the user’s true label. Unlike prior work on strategic classification, in our setting, users *cannot manipulate their data*, but *can select between different services* and *vary their level of usage*. Thus, the features x_i and label y_i of each user i are fixed. Instead, users strategically adjust their *usage*, denoted by $A_{ij} \in \mathbb{R}_+$ for user i and service j . Each user i selects m usage values A_{i1}, \dots, A_{im} . These values are non-negative and real-valued, meaning that the user can modulate the total amount of usage. For example, in the digital credit setting, usage could correspond to the number

or amount of loans.

One component of the strategic user objective is the utility from positive classification. This utility $u : \mathcal{X} \times \mathcal{H} \rightarrow \mathbb{R}$ depends on the user features and the classifier. We make the following assumptions about the utility:

Assumption 1 (User utility). *For any $h_1, h_2 \in \mathcal{H}$ and $x \in \mathcal{X}$ such that $h_1(x) = -1$ and $h_2(x) = 1$, $u(h_1, x) \leq 0 < u(h_2, x)$.*

This assumption states the intuition that users seek positive classification. Furthermore, we allow the utility to depend on the classifier h , not merely the individual classification $h(x)$. This allows users to be sensitive to other qualities of the classifier, such as the margin, i.e. the distance from a negative classification. We illustrate this in the following example:

Example 2. *For a linear model class where classifiers are represented by their weights θ , the 0-1 utility is given by $u(x, \theta) = \mathbf{1}\{\theta^\top \varphi(x) > 0\}$. This models users who care only about whether their classification is positive. The linear utility is given by $u(x, \theta) = \theta^\top \varphi(x)$. This models users who care about their margin.*

The benefit that a user i receives from a service j depends on both utility and usage: $A_{ij}u(x_i, h_j)$.

The second component of the user objective depends only on their total usage: $\sum_{j=1}^m A_{ij}$. In particular, higher total usage corresponds to a lower overall objective value. This models an opportunity cost or a penalty for large usages on a per-user basis. In the digital credit example, individuals are disincentivized from applying for loans too often.

Putting these components together, the overall strategic objective for a user i is to

maximize

$$\sum_{j=1}^m A_{ij} u(x_i, h_j) - \frac{1}{q} \left(\sum_{j=1}^m A_{ij} \right)^q \quad (2.1)$$

where the power $q > 1$ sets the opportunity cost to increase superlinearly, so the marginal utility of usage is non-increasing. To understand this objective, consider the best response of a user i for fixed services h_1, \dots, h_m . If no service offers a positive classification, then by Assumption 1, no utility will be positive, and thus the best response will be zero usage. If there is a service j offering uniquely maximum utility $u(x_i, h_j) > 0$, then the best response is to ignore any other service, i.e., $A_{ik} = 0$ for all $k \neq j$. The best response usage will be $A_{ij} = u(h_j, x_i)^{\frac{1}{1-q}}$, thus depending on the maximal utility in addition to the power q .

If several services offer equal and maximal utility to the user, then the best response is not unique. It includes any configuration of usage distributed over these equivalent services, with fixed total usage. In other words, users consider only the services offering them a sufficiently positive classification, and their usage magnitude is determined by the magnitude of the utility.

We note that implicit in this behavior is that users have full knowledge of service classifiers. This is a common assumption in strategic classification literature and can be motivated through openly released classifiers or hearsay. In the imperfect knowledge setting, we add that services observe users that they classify as negative, making the game easier for the services.

2.2 Learning Updates for Services

Services deploy classifiers trained on data that they have observed. We take the number of services to be much smaller than the number of users $m \ll n$. To each service, the large and fluid user pool is represented as a *distribution* over features and labels. The distribution observed by a service depends on the usages so that the weight on data point (x_i, y_i) for service j is proportional to the usage A_{ij} . Importantly, this means that when the usage is zero, the service has no information about the datapoint. Users who never choose to use a particular service are essentially invisible to the service.

We model the process of training a classifier by expected loss minimization. The loss $\ell : \mathcal{H} \times \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$ quantifies the error of a given classifier on a given data point. We make the following assumption about the loss function.

Assumption 2 (Loss-utility relationship). *For all $h \in \mathcal{H}$, the loss is non-negative, $-y\ell(h, x, y)$ strictly monotonically increases with $u(x, h)$, and there exists $v > 0$ such that $u(x, h) = 0 \implies \ell(h, x, y) = v$.*

For a positive user ($y = +1$), this assumption gives a negative relationship between classification utility and the service’s quantification of error: high utility corresponds to low error, so the goals of users and services are aligned. For a negative user ($y = -1$), it is a positive relationship: high utility corresponds to high error, so the user is at odds with the service. The requirement of the existence of a $v > 0$ meanwhile implies that zero loss gives a user the same sign utility as their label.

Example 3. *For a linear model class where classifiers are represented by their weights θ , the 0-1 loss is $\ell(\theta, x, y) = \mathbf{1}\{\theta^\top \varphi(x) \cdot y > 0\}$, which corresponds to the 0-1 utility. The hinge loss is defined as $\ell(\theta, x, y) = \max\{1 - \theta^\top \varphi(x) \cdot y, 0\}$, which corresponds to the linear utility.*

Recall that as discussed above, realizability means that the model class \mathcal{H} includes a perfect classifier. We now formalize this intuition in terms of the loss function.

Assumption 3 (Realizability). *There exists a classifier $h \in \mathcal{H}$ such that $\ell(h, x_i, y_i) = 0$ for all $i = 1, \dots, n$.*

Example 4. *For the linear model class and either the 0-1 loss or the hinge loss, the realizability assumption is satisfied as long as the features $\{\varphi(x_1), \dots, \varphi(x_n)\}$ are linearly separable with a strictly positive margin.*

This assumption is equivalent to ensuring that there exists a classifier that achieves zero expected loss on the entire population of users. For a given user distribution \mathcal{D} , the expected loss of a classifier h is $\mathbb{E}_{x,y \sim \mathcal{D}}[\ell(h, x, y)]$. When this distribution is defined for a service j based on usages A_{1j}, \dots, A_{nj} , it can be simplified¹ to

$$\sum_{i=1}^n \frac{A_{ij}}{\sum_{k=1}^n A_{kj}} \ell(h, x_i, y_i). \quad (2.2)$$

Therefore, when user distributions are defined by usages, the expected loss depends on the usages as well.

Many works on strategic classification considering the setting where users manipulate their features consider a principle-agent game between a single service and multiple users (13). These works focus on showing the existence of a desirable equilibrium, where desirability is defined as correct classifications with respect to true user labels. In our setting, the realizability assumption implies that such good outcomes are possible. However, because we do not model omniscient services—they observe data depending on strategic usage—arriving at such an equilibrium through interactions requires a nontrivial analysis of dynamical and transient behavior.

¹If $\sum_{k=1}^n A_{kj} = 0$ for some service, we adopt the convention that for all users i , the fraction $A_{ij}/(\sum_{k=1}^n A_{kj}) = 0$.

2.3 Interaction Dynamics

We study the dynamics of repeated interactions between users and services, taking a cue from other lines of work on strategic classification (32), endogenous distribution shift (27), and usage dynamics (6). In these dynamics, users act according to their strategic objective, while services update their classifiers by retraining. We thus index classifiers and usage by time, and denote the state of the dynamics as

$$H^t = (h_1^t, \dots, h_m^t), \quad A^t \in \mathbb{R}^{n \times m}$$

We consider user best response dynamics, so at time t , user i selects a usage to maximize (2.1) given models H^t . We allow users to employ any tie-breaking scheme (even a non-deterministic one) when there is not a unique maximizing model. Consider the following joint update:

$$A^t \in \operatorname{argmax}_{A \in \mathbb{R}_+^{n \times m}} \sum_{i=1}^n \left[\sum_{j=1}^m A_{ij} u(x_i, h_j^t) - \frac{1}{q} \left[\sum_{j=1}^m A_{ij} \right]^q \right] \quad (2.3)$$

We consider this joint update for simplicity of exposition, noting that it is equivalent to any order of independent user updates, due to the separability of the objective.

We consider services that repeatedly retrain their classifiers. The most naive retraining approach minimizes the immediate expected loss (2.2) given usages observed A^t , i.e. over the current user distribution. Such an approach has been studied in the traditional strategic classification setting (32; 27), where it has been shown to converge to desirable equilibria. As we will see in Section 3.1, this *memoryless* retraining is a poor fit for the strategic usage setting.

We therefore consider retraining updates which minimize the weighted sum of prior expected losses; that is, the update to models H^{t+1} considers the expected loss (2.2) induced by A^t, A^{t-1}, \dots . By the linearity of expectation, this is equivalent to minimizing

the expected loss over a *distribution with memory*. Thus, rather than depending on current usage A^t , classifiers are selected according to an average loss depending on some memory $M^t \in \mathbb{R}_+^{n \times m}$. We define memory as follows, using a discount factor $p \geq 0$:

$$M^t = \frac{A^t}{1+p} + \frac{pM^{t-1}}{1+p} \quad (2.4)$$

By convention, $M^0 = 0$. The memoryless case is $p = 0$, where the loss depends only on the current timestep. For $p > 0$, prior observations are retained, but with discounted weight, so that timesteps further in the past contribute less. As p increases, the memory is “stronger”: prior timesteps have more influence on the loss.

Therefore, the retraining dynamics are defined by a service j selecting a classifier to minimize the expected loss defined by the memory M_t . The simultaneous joint update for services is

$$H^{t+1} \in \operatorname{argmin}_{H \in \mathcal{H}^m} \sum_{j=1}^m \sum_{i=1}^n \frac{M_{ij}^t}{\sum_{k=1}^n M_{kj}^t} \ell(h_j, x_i, y_i). \quad (2.5)$$

Due to the separability of the objective over services, this joint update is equivalent to any order of independent service updates. We thus consider the simultaneous update for simplicity of exposition.

We allow classifiers to employ many types of tie-breaking schemes when there is not a uniquely optimal classifier; however, we introduce a requirement of sticky tie-breaking.

Definition 1. *Let there be an update schema where some f^{t+1} is selected to minimize an expected loss as given by $f^{t+1} \in \operatorname{argmin}_{f \in \mathcal{F}} L^t(f)$. The update is called sticky when if for a given f^t , $L^{t-1}(f^t) = L^t(f^t)$, then it holds that $f^{t+1} = f^t$.*

For example, when there is a norm defined on \mathcal{F} , the minimum-norm update rule is

$$f^{t+1} = \operatorname{argmin}_{f \in \mathcal{F}} \|f\|_{\mathcal{F}} \quad \text{s.t.} \quad f \in \operatorname{argmin}_{g \in \mathcal{F}} L^t(g).$$

This update rule satisfies the stickiness property. This requirement is grounded in real-world settings as changing classifiers could often be costly or undesirable, and so when no explicit advantage is derived from replacing a model, it might be more advantageous to retain the old model rather than approve and release a new one.

We consider alternating updates between services and users, where the memory evolves according to (2.4).

Assumption 4. *Given state (H^t, A^t) at time t , the classifier update to H^{t+1} is sticky and satisfies (2.5), and the usage update to A^{t+1} satisfies (2.3).*

This means that first the services update $H^t \rightarrow H^{t+1}$ based on current usage A^t , and then the users best-respond $A^t \rightarrow A^{t+1}$ based on the updated services H^{t+1} . We remark that our analysis techniques could extend to a round-robin of updates involving varied sequences of service and user updates. We include a further discussion in the appendix and focus on the joint updates here for simplicity of exposition and notation.

CHAPTER 3

THEORETICAL RESULTS

In this section, we formally introduce our theoretical results concerning the dynamics and convergence of the interaction defined by Assumption 4. We will show that under mild assumptions when the memory is non-zero, the system will reach a desirable point within a finite number of steps. Here, we define desirability from the perspective of accurate classification.

Definition 2. A state (H, A) is zero-loss if all services j satisfy: 1) $A_{ij}\ell(h_j, x_i, y_i) = 0$ for all i and 2) $u(x_i, h_j) \leq 0$ for all i with $y_i = -1$.

Zero-loss states are desirable from many perspectives. The first condition means that all users with nonzero usage have minimal loss. This means that all services make accurate classifications on the populations they observe. For positive users, it means maximal utility and correctly positive classifications. The second condition means that negative users will receive zero utility and will thus choose not to engage in any service.

Beyond desirability from the perspective of utility and loss, zero-loss points play an important role in understanding the convergence of the interaction dynamics. In Section 3.2, we show that for memory parameter $p > 0$, when a state is zero-loss, so are all future states. Thus the zero-loss property defines a set which is *invariant* under the service retraining and user best response updates. In Section 3.3, we further show that the dynamics will reach a zero-loss state in finite time.

All results are presented under Assumptions 1, 2 3, and 4.




Figure 3.1: Five datapoint example described in the proof of Proposition 1. Negative points are represented by $-$ and positive points by $+$, where boldface indicates two overlapping points. The dashed and solid lines represent the oscillating classifiers, with the dotted line representing a zero-loss classifier.

3.1 Importance of Memory

While naive repeated retraining can be a successful strategy for strategic classification problems in the context of data manipulation (32), in this section we show that it can catastrophically fail.

Proposition 1. *In the memoryless $p = 0$ setting, there exist initializations in which the state (H, A) never converges.*

We prove this result by providing an example that leads to oscillations. We construct a dataset and an initial configuration satisfying all assumptions and show that classifiers and usages will fluctuate indefinitely. The dataset is illustrated in Figure 3.1.

Proof Sketch. We consider a setting of five users and two classifiers, illustrated in Figure 3.1 and described in detail in Section 4.2. We show that for $p = 0$, there are initial conditions that lead to perpendicular oscillating classifiers (dashed and solid in the fig-

ure) rather than zero-loss (for example, dotted). \square

In contrast to memoryless updates, when $p > 0$, services accumulate knowledge about data distribution over timesteps. We will show that in our setting, nonzero memory is enough to guarantee convergence to a zero-loss state. Towards that goal, we first make precise the ability of a service to accumulate knowledge.

Lemma 2. *For every user service pair i, j such that $M_{ij}^t > 0$, $\ell(h_j^{t+1}, x_i, y_i) = 0$. Therefore, when $p > 0$, if $A_{ij}^t > 0$ for any timestep t , then for all further timesteps $\tau > t$ it must hold that $\ell(h_j^\tau, x_i, y_i) = 0$.*

The proof of this result, and the missing proofs of all remaining results, are presented in the appendix.

3.2 Characterizing Invariance

So far, we have established that, in order to avoid oscillations, services must have memory $p > 0$ when retraining their classifiers. Now, we turn to questions of convergence. However, unlike many related works, we do not study a setting that necessarily admits fixed points, much less unique fixed points. This occurs because we allow arbitrary user best response updates and assume realizability. This means that it is plausible for users to vary their usage among several different services indefinitely, so long as those users are positive and the services classify them correctly. As a result, some care is required to define the appropriate notion of “convergence.” Instead of arguing about fixed points, we turn our attention to an *invariant set* defined by the zero-loss property. The following proposition formalizes the idea of invariance for zero-loss points.

Proposition 3. *If a state (H^t, A^t) is zero-loss at time t , then all future states are zero-loss*

for all times $\tau \geq t$.

The proof of this proposition follows from the sticky assumption and users being unable to impart loss on the services without engaging in suboptimal actions. We further establish that the sticky assumption is necessary for classifier tie-breaking in order to guarantee this invariance.

Proposition 4. *Without sticky tie-breaking (Assumption 4), regardless of the value of p , the existence of a zero-loss equilibrium (H^t, A^t) at a timestep t does not guarantee the existence of zero-loss equilibria for further timesteps $\tau > t$.*

Proof Sketch. We consider a setting where a service is initialized with a model only giving positive usage to positive points, with other negative points never seen. This may be observed to be a zero-loss state. At the next timestep, or some future one, the service may feasibly re-sample a new model that gives positive usage to the negative points as they were never observed, resulting in a new state where the zero-loss conditions are violated. \square

3.3 Convergence

Finally, we turn to *convergence*: regardless of the initial state of classifiers and users, we show that the interaction dynamics will lead to a zero-loss point within finite time.

Over the course of interactions between services and users, the services collect data. Each time a new user selects a service, the service may respond by updating its classifier— in particular, a change must occur if the new user is incorrectly classified. Towards ensuring convergence to a zero-loss point, we hope to show that this process

terminates. The following lemma presents a sufficient condition for reaching a zero-loss point in terms of the usage update.

Lemma 5. *For any timestep t if there exists no values $M_{i,j}^{t-1} = 0$ such that $A_{i,j}^t > 0$, then (H^t, A^t) is zero-loss.*

The proof of this follows from extending Lemma 2 to the following timestep.

With this lemma in hand, we are ready to prove the main result, which shows that services and users will converge to a zero-loss equilibrium in finitely many steps.

Theorem 6. *Given nonzero memory $p > 0$, there is a finite time $t \in \mathbb{N}$ after which for all $\tau > t$, (H^τ, A^τ) is zero-loss.*

Proof Sketch. We first argue that there are only finitely many timesteps in which the condition in Lemma 5 can occur. Therefore, the system must reach a zero-loss point, and by Proposition 3, the classifiers and usages must continue to be zero-loss. \square

CHAPTER 4

EXPERIMENTS

4.1 EXPERIMENTS

We illustrate our results with experiments of simulated strategic usage behavior. We first consider a synthetic data setting similar to that found in the proof of Proposition 1 to illustrate the importance of memory. We then perform experiments on real data in the setting of the Banknote Authentication task (21) and the realistic synthetic data of the Bank Account Fraud task (18). These experiments showcase the importance of memory and illustrate the complex and initialization-dependent dynamical behaviors that strategic usage dynamics can induce.

4.2 Five Points Dataset

We begin with a synthetic example with only five points to illustrate oscillation in non-memory cases, with features:

$$\{(1, 1), (1, 1), (-1, 1), (1, -1), (-1, -1)\}$$

and labels $\{1, 1, -1, -1, -1\}$ respectively. This dataset is clearly separable by linear classifiers with $\varphi(x) = [x, 1]$ and $d = 3$ (Example 1). We consider users acting strategically according to the linear utility defined in Example 2 with services updating according to the hinge loss defined in Example 3. We consider $m = 2$ services with initial models $\theta = [1, 0, 0]$ and $\theta = [0, 1, 0]$; the dividing hyperplanes are perpendicular to one another, both giving positive classifications to the coincident positive points, but while model 0 gives positive classification to $(1, -1)$ the other gives positive classification to $(-1, 1)$. It

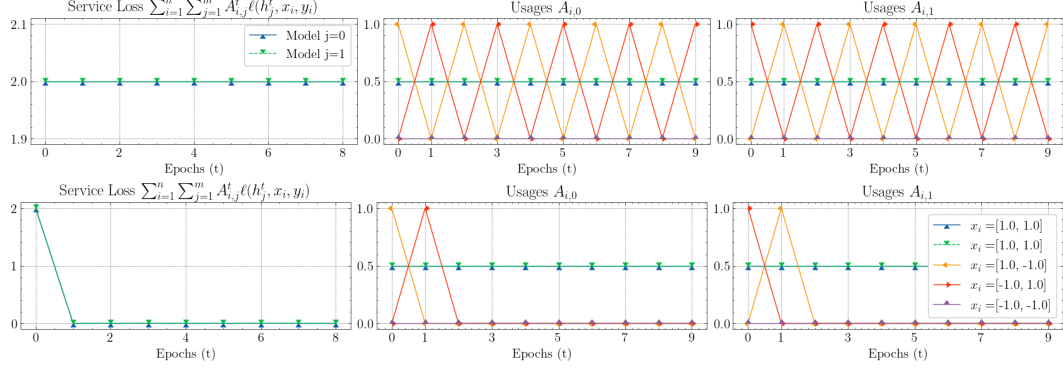


Figure 4.1: 5-Points dataset; the top three graphs give the $p = 0$ case while the bottom three give $p = 0.5$. Service loss is calculated after the user update but before the service update, and usages are displayed for each of the five points with the middle graphs giving the usages for model $j = 0$ and the right graphs giving the usages for model $j = 1$.

may be observed that this setting is identical to that of the proof of Proposition 1 and is illustrated in Figure 3.1. Note that the fifth user is a negative point who will never choose to use either classifier, and thus will not be seen by them.

For tie-breaking, services choose models by minimizing the norm of θ (subject to achieving zero loss), while users split usage equally between models that assign equal utility to them. We run two experiments with this synthetic dataset: one being memoryless with $p = 0$ and the other using $p = 0.5$. This demonstrates how the memoryless setting may lead to oscillations while the inclusion of memory ensures convergence.

Results are presented in Figure 4.1, which plots the loss and usage of each service. In the $p = 0$ case, there is clear oscillation in the usages—the users at $(1, -1)$ and $(-1, 1)$ alternate between the two models. In contrast, for $p = 0.5$, the usages converge after the second epoch. Perhaps more illuminatingly the loss drops to zero: this indicates that the services have converged to a zero-loss point (Definition 2) and no longer change between updates, unlike in the $p = 0$ case where classifiers swap directions each timestep. It must also be noted that the usages of the negative users specifically converge to 0 in the $p > 0$ case. Since the services correctly learn to assign negative classifications to the negative

users, even the unseen users at $(-1, -1)$, there is no incentive for nonzero usage from negative users.

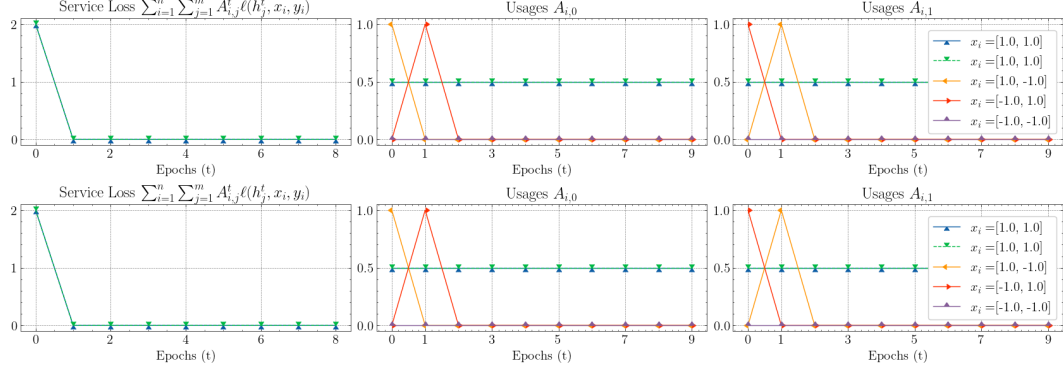


Figure 4.2: 5-Points dataset; the top three graphs give the $p = 0.1$ case while the bottom three give $p = 1.0$.

In order to explore the effect of different values of $p > 0$ on convergence, an ablation on p is presented in Figure 4.2. The top set of plots gives service loss and usages for $p = 0.1$ and the bottom set gives the same for $p = 1$; however, both values of p give the same usages and losses $p = 0.5$ across epochs. This gives us that only the cases of $p = 0$ and $p > 0$ produce a meaningful difference empirically.

4.3 Banknote Authentication

We instantiate a semi-synthetic simulation with real-world data coming from the Banknote Authentication dataset (21). This dataset involves a binary classification problem to detect whether a banknote is genuine or forged. Banks are modeled as services and update their forgery-detection classifiers in order to reject forged banknotes while accepting genuine ones. Meanwhile, users are individuals who seek, for practical purposes, banks that allow them to cash in their notes. At the same time, banks update their forgery detection models to keep up with trends in the forgery industry. We re-

mark that our *strategic usage* setting corresponds to the short-term dynamics of banks becoming aware of forged bank notes that are in circulation, while the classical feature manipulation setting corresponds to innovations in forgery techniques.

Features are derived from images of banknote-like specimens; specifically, they are extracted using a wavelet transform tool, resulting in $\mathcal{X} = \mathbb{R}^4$. Each sample additionally comes with a binary label. We apply the following preprocessing: we normalize the mean and variance of the features, and we transform the labels $\{+1, -1\}$.

We simulate services using scikit-learn support vector classification (SVC) (26) with a radial basis function (RBF) kernel using radius $\gamma = 1$ and regularization parameter $C = 10^{10}$. This setting corresponds to linear models (Example 1) with an infinite dimensional feature function φ trained with hinge loss (Example 3). The large value of C corresponds to a small regularization weight, which approximates simply selecting the minimum norm model among all loss minimizers. In the memoryless $p = 0$ setting, services have no negative users in their loss objective at various points in time. This violates the preconditions for the scikit-learn SVC fit function, so in these instances, we preserve weights from the previous timestep.

We initialize this setting by revealing one positive and one negative user at random to each service at a timestep $t = -1$ in order to train models h_j^0 for all $j \in \{1, \dots, m\}$. Random seeds for choosing these users are held constant at 100 between runs for consistency, and users tie-break through splitting usage evenly between services that provide equal usage to them.

Figure 4.3 presents simulation results for $m = 5$ services, plotting the total usage of the positive and negative class over time (i.e., of genuine and forged banknotes in circulation). In the memoryless $p = 0$ case, we observe some complex transients until

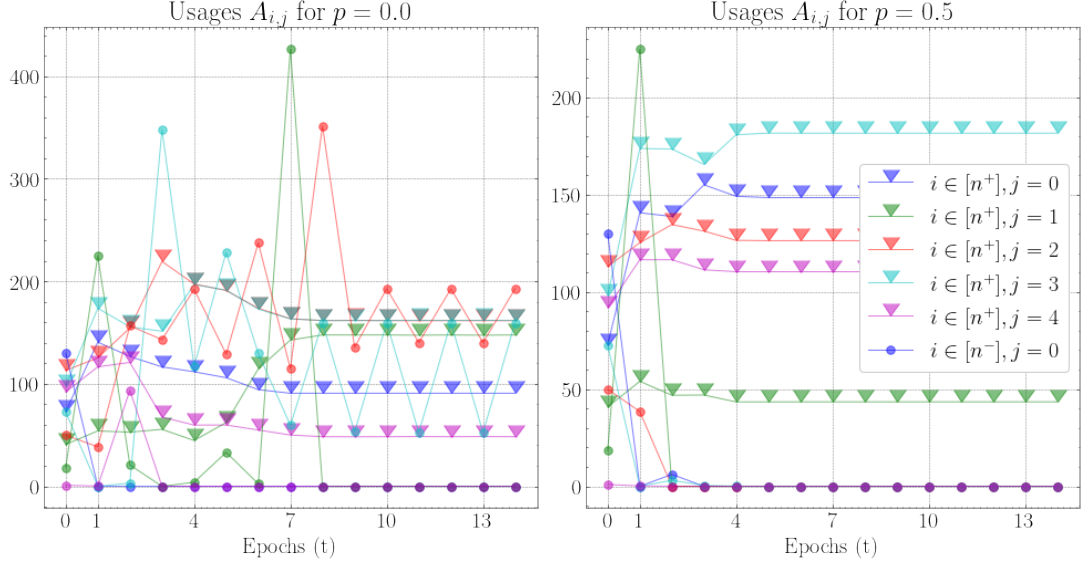


Figure 4.3: Banknote Authentication dataset; each graph gives the positive and negative usages of each of the five models; triangle markers above the lines indicate positive usage while below the lines indicate negative, with colors giving which model the line refers to. The left graph gives the no-memory $p = 0$ setting, while the graph on the right gives the $p > 0$ setting. Model order, and hence their colors, are meaningless due to the random initialization.

the tenth timestep, after which point we see oscillation between models $j = 2$ and $j = 3$. The remaining models stopped observing negative points and henceforth stopped updating; however, it is important to note that which models converged and which models continued oscillating is initialization dependent: changing the seed resulted in different dynamics, including changing the values that models' usages converged or oscillated around, in addition to changing the numbers of models converging versus oscillating.

In the nonzero memory case $p > 0$, we once again observe convergence to a zero-loss point, this time after four timesteps—the longer transient period resulting from the larger number of services that users can choose between. At this equilibrium, each model correctly assigns negative classifications to all the negative users, while the positive users are divided between the services.

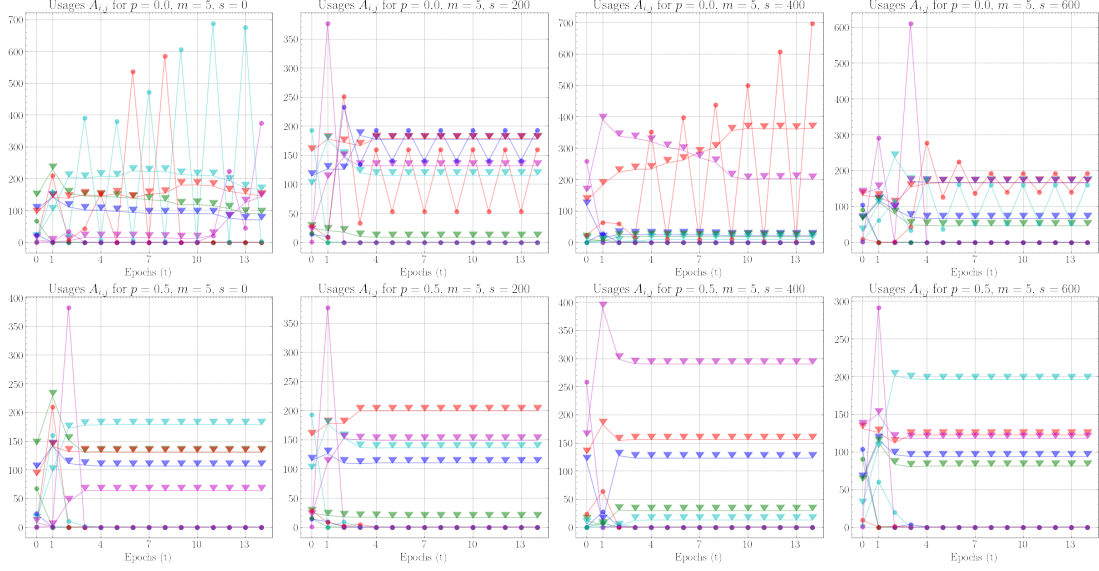


Figure 4.4: Banknote Authentication dataset; ablating on the seed (s). The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.

The particular configuration of positive users is an arbitrary result of the random initialization. We illustrate the potential for a variety of outcomes depending on the initial seed in Figure 4.4. Generally, this shows that the initial conditions can have drastic effects on both the time to convergence and the final stable state.

Further interrogating the effect of hyperparameters on convergence, ablation on the number of services m is presented in Figure 4.5. Plots for $m = 1, 2, 3, 4$ are given, both in the zero memory and nonzero memory cases. This illustrates that as the number of services increases, convergence may take longer due to services interfering with each other and disincentivizing users to reveal themselves to other services through usage. Note that due to the static seed, between graphs models are shown the same initial users when the model is present.

Finally, to inspect the effect of cost power factor q on convergence, we provide

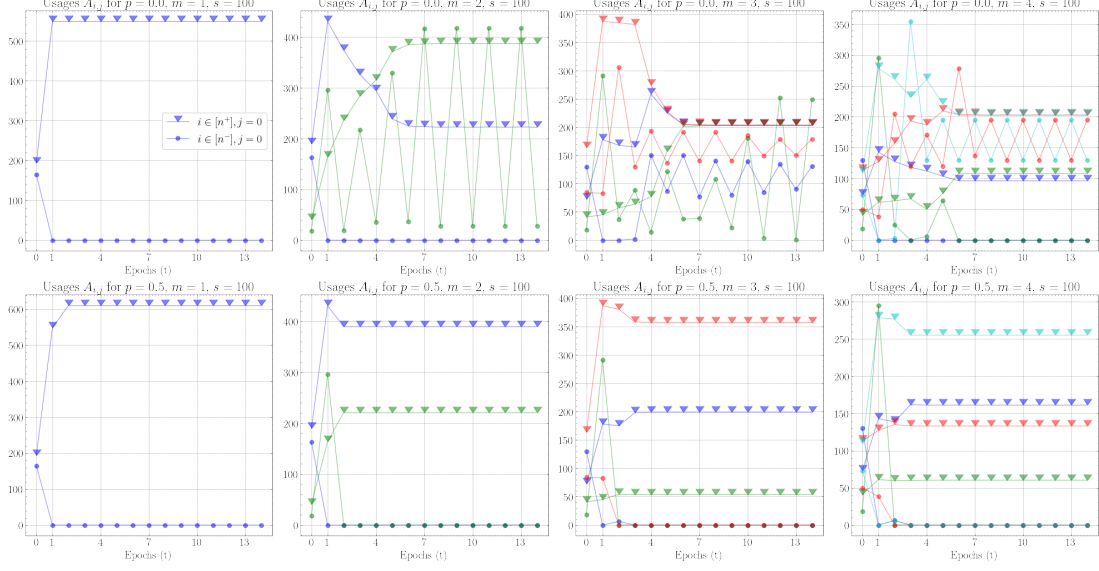


Figure 4.5: Banknote Authentication dataset; ablating on m . The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.

usages over epochs for varying values of q both with and without memory in Figure 4.6.

No meaningful relationship with convergence is found for $q > 1$.

4.4 Bank Account Fraud

We include a third dataset as the Bank Account Fraud (BAF) dataset (18). This dataset models the binary classification problem of predicting if a bank account opening is fraudulent or not. We interpret banks as services providing bank accounts, with users being individuals attempting to open accounts. As with the banknote authentication setting, this strategic usage setting would provide the precursor to potential feature manipulation which might occur once fraudulent users realize no bank will provide them with utility.

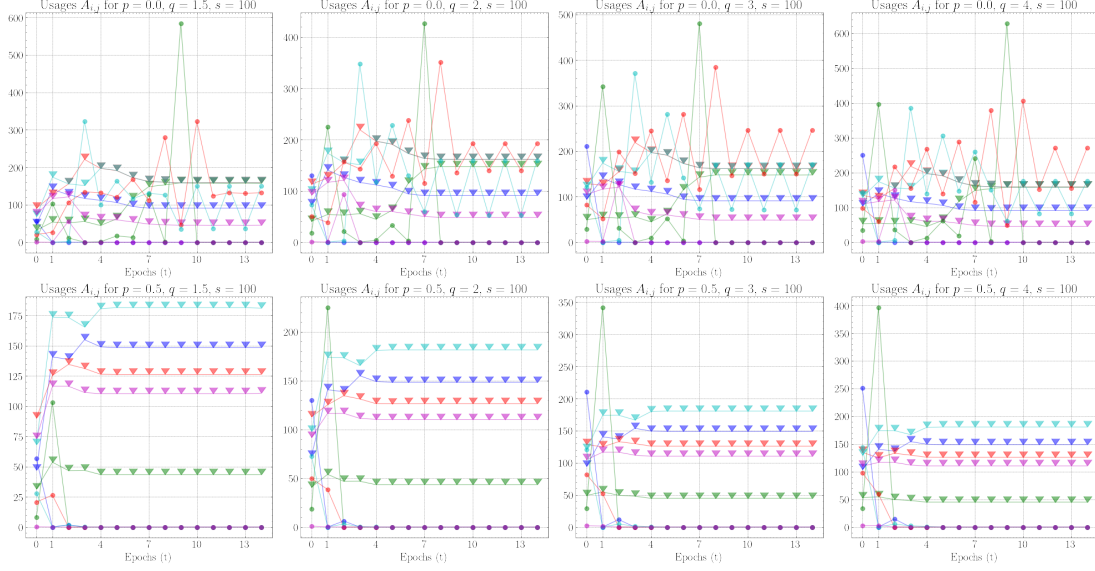


Figure 4.6: Banknote Authentication dataset; ablating on usage cost power factor q . The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.

This dataset suite was designed to stress test the performance and fairness of machine learning models on realistic tabular data generated from a real-world bank account opening fraud dataset. None of the datasets fully satisfy our realizability assumption; we selected Variant V because it was the most separable. We filtered the data points to ensure realizability by fitting a soft-margin support vector machine and removing misclassified points. Each sample has a binary label giving whether the account opening is fraudulent or not, and features include values such as income, customer age, and payment type, exhibiting both real-valued and enumerated types. We transform non-numerical features into one-hot representations before normalizing the mean and variance of all features, and labels are transformed to $\{+1, -1\}$. All experiments were run for $m = 5$ services, with cost power factor $q = 2$. From the realizable subset of the data, we select the first 5,000 positive and 5,000 negative points to speed up runtime.

We use Scikit-learn SVC (26) to simulate services, with a linear kernel and regular-

ization parameter $C = 10^{10}$. Services are initialized in a similar manner to the banknote authentication dataset, and a random seed of 100 is held constant between runs for consistency.




Figure 4.7: Bank Account Fraud dataset; each graph gives the positive and negative usages of each of the five models. Notation is the same as that for Figure 4.3.

Experiments for the $p = 0$ and $p > 0$ can be seen in Figure 4.7, for $m = 5$ services. We once again observe oscillation and non-zero usage for negative users in the memoryless $p = 0$ case, exhibiting the failure of the system to converge to a zero-loss state. The opposite can be seen in the $p > 0$ case, as no change in usage may be observed after the seventh epoch, and all negative users are observed to have left the system.

Results under a variety of initial seeds can be seen in Figure 4.8, including both results with and without memory. These results further corroborate the earlier theoretical findings, demonstrating convergence in the memory case with oscillation commonly occurring in the naïve retraining setting.

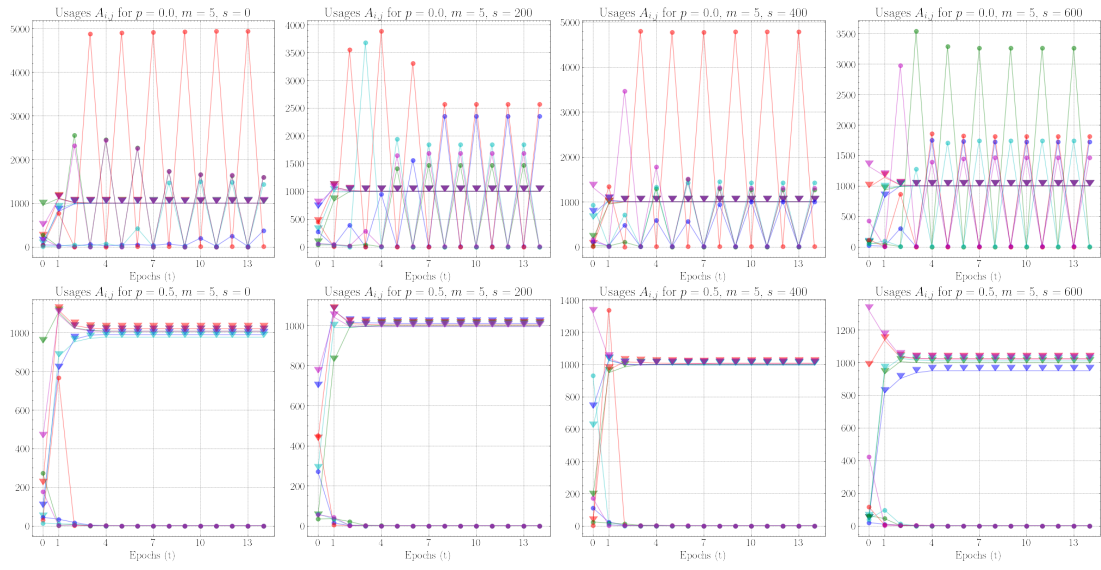


Figure 4.8: Bank Account Fraud dataset; ablating on the seed (s). The top four graphs give the $p = 0$ case while the bottom four give $p = 0.5$. In each graph, triangle markers indicate positive usages while circular markers indicate negative usages; colors indicate the different services.

CHAPTER 5

RELAXING ASSUMPTIONS

In previous chapters, we formalized the strategic usage setting and established a set of sufficient assumptions to guarantee finite-time convergence to a zero-loss state. In this chapter, we shall prod at some of these assumptions and analyze the results of their relaxation on convergence.

5.1 Round Robin Updates

It might not be realistic that users and services update on the same schedule, as is taken in Assumption 4: while prior proofs assume that services conduct one synchronous update based on current usage and users synchronously best respond once based on the updated services at every timestep, it could be that services only update every several years while users may reallocate usage yearly. Additionally, users and services don't necessarily update synchronously, and one might conduct several updates at a time while others only do one. In this setting, a timestep stops being a feasible metric of progression; instead, we generalize to the concept of a round.

Instead of there being one joint user and one joint service update as in a timestep, we generalize rounds to users and services updating asynchronously and differing numbers of times. We maintain three constraints on this system. First, rounds are divided into alternating user update periods and service update periods, such that only users or services are updating at a time. Second, each round must contain at least one user period and one service period. Finally, each user and each service undergoes at least one best response update in each period.

Definition 3. An update step is a time concept in which for some step index t , either a user or service update occurs. A user update step for some user i is defined as a step-index t such that the following occurs:

$$A_i^t \in \operatorname{argmax}_{A_i \in \mathbb{R}_+^m} \sum_{j=1}^m A_{ij} u(x_i, h_j^t) - \frac{1}{q} \left[\sum_{j=1}^m A_{ij} \right]^q$$

$$A_{-i}^t = A_{-i}^{t-1} \quad M^t = M^{t-1} \quad H^t = H^{t-1}$$

A service update step for some service j is defined as a step-index t such that the following occurs:

$$A^t = A^{t-1} \quad M_j^{t\top} = \frac{A_j^{t\top}}{1+p} + \frac{pM_j^{t-1\top}}{1+p} \quad M_{-j}^{t\top} = M_{-j}^{t-1\top}$$

$$h_j^t \in \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^n \frac{M_{ij}^t}{\sum_{k=1}^n M_{kj}^t} \ell(h_j, x_i, y_i) \quad H_{-j}^t = H_{-j}^{t-1}$$

Definition 4. A round is an ordered tuple of a user round and a service round $r = ([u_t], [s_t])$. The user round is a set of consecutive update steps u_t satisfying for all users i , there exists some step t_i such that user i underwent an update. The service round is a set of consecutive update steps s_t satisfying for all services j , there exists some step t_j such that service j underwent an update.

We note that though we require all users to have undergone an update in a user round, this doesn't disclude the possibility of service updates happening during the user round or of users undergoing multiple updates in a single user round. This is simply an encapsulating concept in order to ensure some metric of progression: the same goes for a service round. With this concept of a round, we may achieve a similar result to that found in the synchronous update setting.

Proposition 7. Given nonzero memory $p > 0$, there there are a finite number of rounds (Definition 4) $r \in \mathbb{N}$ after which for all $\rho > r$, (H^ρ, A^ρ) is zero-loss.

5.2 Non-Realizability

While previous results analyzed a setting involving a realizability assumption (Assumption 3), ultimately there do exist a wide variety of scenarios involving a non-realizable domain. Relaxation of this assumption is furthermore necessary for an interesting analysis of a variety of extended settings, including stochastic user labels and users hiding labels. As such, in this section, we take the opportunity to explore the relaxation of the realizability assumption and its effects on convergence.

It is immediately obvious that the concept of a zero-loss state is no longer achievable in this setting, breaking Proposition 3. In order to define a meaningful result despite the non-existence of perfect classifiers, let us use the concept of a stable state.

Definition 5. A state (H^τ, A^τ) is a stable state if for all $t > \tau$, $H^t = H^\tau$.

In this stable state, we allow for users to reassign usage between services in equivalence classes; however, if the loss achieved by these users doesn't change in between timesteps, we expect the service models H^t to remain constant.

With this new definition, we may first observe that in the unrealizable setting, even this stable state may be unachievable due to issues with memory oscillation. This is formalized below.

Proposition 8. When there doesn't exist any $h \in H$ such that $\ell(h, x_i, y_i) = 0$ for all $i = 1, \dots, n$, there exist initializations in which the state (H, A) never reaches a stable state.

Proof Sketch. We consider a setting of a positive user, a negative user, and one linear classifier with a margin. Suppose the two users are of a distance less than the classifier's margin away from one another. In that case, the memory updates will regularly cause

the classifier to ignore the negative point and classify it as positive in order to make the positive point satisfy its margin. \square

In order to handle this issue of oscillation along the margins, we modify our definition of memory. Instead of weighted memory updates as in Equation 2.4, we take advantage of the static finite dataset and treat memory as a 0 – 1 function of unseen and seen. This setting can be considered as services memorizing individual users, rather than learning a distribution over time.

$$M^t = \begin{cases} A^t & \text{if } p = 0 \\ \mathbf{1}\{\max(A^t, M^{t-1})\} & \text{if } p > 0 \end{cases} \quad (5.1)$$

We conduct empirical experiments on this setting using the Bank Account Fraud dataset, and results are presented below. The initialization is done with a random seed of 100, to match the earlier results.




Figure 5.1: Bank Account Fraud dataset in the unrealizable setting. The left plot gives the $p = 0$ case, the middle plot gives the $p = 0.5$ case using Equation 2.4 for the memory update, and the right plot gives the $p = 0.5$ case using Equation 5.1 for the memory update.

In Figure 5.1, we first observe Equation 2.4 and Equation 5.1 are equivalent for the $p = 0$ case, as is reflected in the leftmost plot where normal oscillation is observed.

In the middle plot, giving the results of the $p > 0$ case when acted upon by memory updating under Equation 2.4 empirically we may observe convergence to a narrower class of solutions; however, there is no achievement of a stable state. This is likely due to negative points periodically losing weight in memory and being ignored until they are observed again. Finally, the rightmost plot shows a stable state after the fifth epoch in the $p > 0$ case using Equation 5.1. As memory no longer updates once a stable classifier is found by each user, we could expect such results to hold theoretically as well.

5.3 Sampled Users

A final meaningful relaxation we consider is that of users not being a finite static dataset, but rather sampled from a distribution. This would mimic the real-world setting of users entering and leaving the system over time, and make meaningful the implementation of memory as a weighted sum over time.

In this setting, it is immediately obvious that finite timesteps cannot be our metric of consideration, as if users are sampled stochastically, it could be that new users provide no new information in a given timestep, by being classified as negative and as such not being seen by the services. As such, we shall analyze convergence with respect to a mistake bound. Furthermore, due to the stochastic nature of users being revealed, Proposition 3 cannot be guaranteed as a new user may be revealed that makes a zero-loss state from a previous state be violated in the next: as such, we must once again revert to the stable state as defined earlier.

If \mathcal{H} is unconstrained, it still can necessarily take an infinite number of mistakes to reach a stable state, as \mathcal{H} could be a tabular model class where infinite mistakes would be required to cover a continuous feature space. We conduct empirical experiments

on this setting using the Banknote Authentication dataset, with results being presented below.

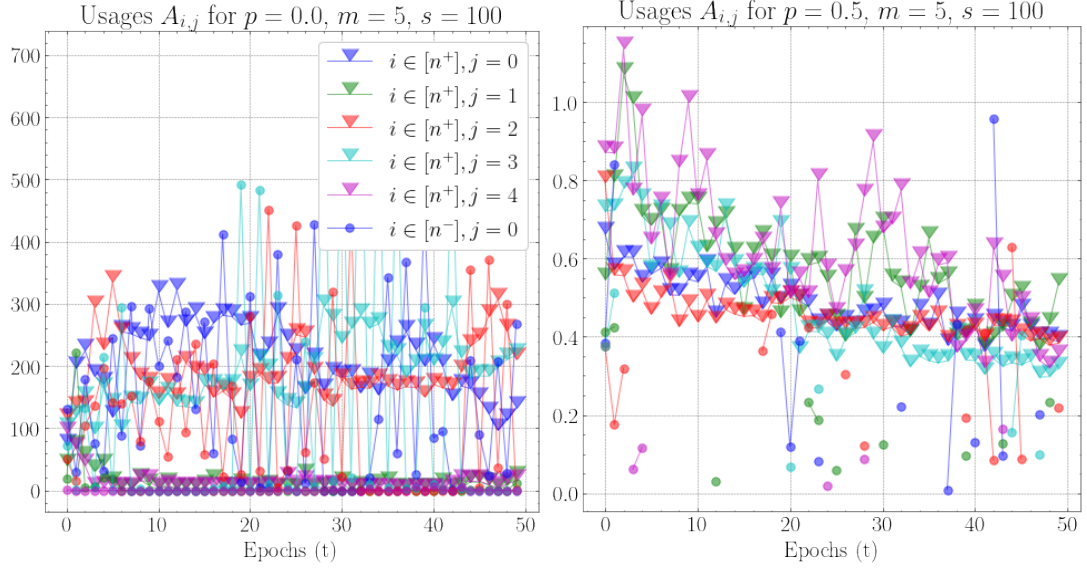


Figure 5.2: Banknote Authentication dataset in the sampled users setting.

Results depicting the effect of sampled users on convergence can be observed in Figure 5.2. The random seed is set to 100 in line with prior experiments, and at every timestep, 100 users are sampled and reveal their usages to five services. In the $p = 0$ case, we observe a lack of convergence as is standard with earlier experiments; however, a slight decrease in variance over time of positive points might be observable for the $p > 0$ case. This trend indicates that if further epochs were run, we might eventually observe an achievement of a stable state; however, there are no theoretical guarantees for when this might be.

CHAPTER 6

CONCLUSION AND DISCUSSION

This work focuses on interaction dynamics between strategic users and multiple learners in an interactive setting. We formalize *strategic usage*, in which individuals choose to use services strategically in order to pursue a positive classification. Since usage presents their data to said services, the latter respond by retraining their classifiers, to the potential benefit or detriment of these strategic individuals. We provide sufficient conditions to guarantee a finite-time equilibrium, including the addition of memory to classifier retraining updates. Our work is both an extension of and in contrast to the strategic classification setting, which primarily explores dynamics involving strategic user feature modification in pursuit of goals. We remark that while several works have raised concern as to the adverse social outcomes entailed by strategic classification (23; 4; 16) our setting ensures that all true positives may receive positive utility at equilibrium.

As the first to study the setting of strategic usage, our work raises several areas for future extension. Though the realizable setting is well motivated for modern ML, it is natural to analyze the setting in which no single classifier can correctly classify the entire dataset. Another natural extension is to consider explicit competition between services, rather than retraining blindly, unaware of the existence of models being deployed by other services in the system. Similarly, we study un-coordinated and short-term user objectives, but an exploration of long-term strategic planning and optimization might give insight into additional real-world phenomena. Finally, it would be interesting to consider the relationship between a finite user pool and an underlying population-level distribution.

APPENDIX A

MISSING PROOFS

In this section, we restate the main theoretical results, include their full proofs, and introduce a handful of auxiliary results.

A.1 Oscillations and Counter-Examples

In this section, we present examples of dynamics that do not converge or that are not zero-loss invariant.

Proposition 1. *In the memoryless $p = 0$ setting, there exist initializations in which the state (H, A) never converges.*

Proof of Proposition 1. Let there exist a set of users with features $X = \{(1, 1), (1, 1), (-1, 1), (1, -1), (-1, -1)\}$ and labels $Y = \{1, 1, -1, -1, -1\}$, selecting between two services where both services choose models from the linear model class (Example 1) using feature transformation $\phi(x) = (x, 1)$, using linear utility (Example 2) and hinge loss (Examples 3). Suppose that the initial models are defined through parameters $\theta_1 = [1, 0, 0]^\top$ and $\theta_2 = [0, 1, 0]^\top$, that retraining updates tie-break by choosing the minimum norm classifier, and that users tie-break by dividing usage equally between services.

Optimal updates may be calculated by analyzing stable points of repeated gradient updates on the best-response updates found in Equations 2.3 and 2.5 noting that with

$p = 0$, $M^t = A^t$ for any timestep t . This algebra gives us the following update steps:

$$\begin{aligned}\theta^0 &= \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}^\top & A^0 &= \begin{bmatrix} 0.5 & 0.5 & 1 & 0 & 0 \\ 0.5 & 0.5 & 0 & 1 & 0 \end{bmatrix}^\top \\ \theta^1 &= \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix}^\top & A^1 &= \begin{bmatrix} 0.5 & 0.5 & 0 & 1 & 0 \\ 0.5 & 0.5 & 1 & 0 & 0 \end{bmatrix}^\top \\ \theta^2 &= \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}^\top & A^2 &= \begin{bmatrix} 0.5 & 0.5 & 1 & 0 & 0 \\ 0.5 & 0.5 & 0 & 1 & 0 \end{bmatrix}^\top\end{aligned}$$

It may be observed that $\theta^2 = \theta^0 \neq \theta^1$, and $A^2 = A^0 \neq A^1$. Given the memoryless setting in which (H^{t+1}, A^{t+1}) depend only on the previous (H^t, A^t) , we may conclude that this system will oscillate for all further timesteps and will not reach a fixed state. \square

Proposition 4. *Without sticky tie-breaking (Assumption 4), regardless of the value of p , the existence of a zero-loss equilibrium (H^t, A^t) at a timestep t does not guarantee the existence of zero-loss equilibria for further timesteps $\tau > t$.*

Proof of Proposition 4. Let there exist a set of users with features $X = \{(5), (-5), (0)\}$ and labels $Y = \{1, -1, 1\}$, making usage decisions in a system with one service choosing models from the linear model class (Example 1) using feature transformation $\phi(x) = (x, 1)$, using linear utility (Example 2) and hinge loss (Examples 3). Suppose that the model is defined through parameters $\theta = [1, -1]^\top$, with retraining updates tie-breaking stochastically.

Optimal updates may be calculated by analyzing stable points of repeated gradient updates on the best-response updates found in Equations 2.3 and 2.5. The first update step can be seen as follows:

$$\theta^0 = \begin{bmatrix} 1 & -1 \end{bmatrix}^\top \quad A^0 = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix}^\top$$

It may be noted that this constitutes a zero-loss equilibrium, as for all users i , $A_i^0 \ell(h^0, x_i, y_i) = 0$, and for the negative point, $\ell(h^0, x_i, y_i) < 1$. Despite this, it is feasible for the following second update step to occur due to H not being constrained by sticky tie-breaking:

$$\theta^1 = \begin{bmatrix} 1 & -0.5 \end{bmatrix}^\top \quad A^1 = \begin{bmatrix} 1 & 0 & 0.5 \end{bmatrix}^\top$$

Here, it can be seen that for the point (0), $A_i^1 \ell(h^1, x_i, y_i) = 0.25$. As such, state (H^1, A^1) does not constitute a zero-loss equilibrium, proving the claim. \square

A.2 Invariance

We next turn to our results about the invariance of zero-loss points. We begin with lemmas characterizing important properties of the service retraining update.

Lemma 2. *For every user service pair i, j such that $M_{ij}^t > 0$, $\ell(h_j^{t+1}, x_i, y_i) = 0$. Therefore, when $p > 0$, if $A_{ij}^t > 0$ for any timestep t , then for all further timesteps $\tau > t$ it must hold that $\ell(h_j^\tau, x_i, y_i) = 0$.*

Proof of Lemma 2. By the separability of Equation 2.5 on services, we may analyze the best-response update of a particular service j . Let us call H_j^{t+1} the set of models such that the best-response update $h_j^{t+1} \in H_j^{t+1}$.

$$H_j^{t+1} := \operatorname{argmin}_{h \in \mathcal{H}} \sum_{i=1}^n \frac{M_{ij}^t}{\sum_{k=1}^n M_{kj}^t} \ell(h_j, x_i, y_i)$$

Since terms of the sum where $M_{i,j}^t = 0$ will simply be zero we may simplify the sum, and furthermore in the interest of brevity let us represent $\frac{M_{ij}^t}{\sum_{k=1}^n M_{kj}^t}$ as $r_{i,j}(M^t)$.

$$= \sum_{i \in \{i | M_{i,j}^t > 0\}}^n r_{i,j}(M^t) \ell(h_j, x_i, y_i) \quad (\text{A.1})$$

By Assumption 3, we have that there exists an $h^* \in \mathcal{H}$ such that $\forall i \in \{1, \dots, n\}$, $\ell(h^*, x_i, y_i) = 0$. By the non-negativity of ℓ and M it must hold that $r_{i,j}(M^t)\ell(h_j, x_i, y_i) \geq 0$ for all i, j pairs, and since substituting h^* into expression A.1 would return a sum of 0, it must be that 0 is the minimum value of the sum. For all $h \in \mathcal{H}$ such that $\ell(h, x_i, y_i) > 0$ for some user i where $M_{i,j}^t > 0$, the value of the sum would be greater than 0 and therefore h wouldn't minimize the sum so $h \notin H_j^{t+1}$. Finally, we may conclude that for all $h_j \in H_j^{t+1}$, $\ell(h, x_i, y_i) \leq 0$ for all i such that $M_{i,j}^t > 0$.

For the second statement, we may first observe that at any timestep t , if $A_{i,j}^t > 0$ then by Equation 2.4 and the non-negativity of M , $M_{i,j}^t > 0$. For the $p > 0$ case, if at some timestep t we have that for some user i and service j , $M_{i,j}^t > 0$, then it will hold that $M_{i,j}^{\tau-1} > 0$ for all timesteps $\tau > t$. This may be trivially observed through the non-negativity of A and the memory update given by Equation 2.4.

By the conclusions drawn above, since $M_{i,j}^{\tau-1} > 0$ we may conclude that $\ell(h_j^\tau, x_i, y_i) \leq 0$. □

Next, we introduce an auxiliary lemma that formalizes the best response behavior of users. This lemma makes rigorous the informal discussion at the end of Section 2.1.

Lemma 9. *For a given user i , let there be a set of services J such that $u_0 = u(x_i, h_j^t) > u(x_i, h_k^t)$ for all services $j \in J$ and other services $k \notin J$. The user best response at time t must satisfy $u_0^{1/(q-1)} = \sum_{j \in J} A_{i,j}^t > \sum_{k \notin J} A_{i,k}^t = 0$.*

Proof of Lemma 9. Analyzing the strategic objective for a user i (2.1) and denoting it as L_i^u for convenience, we can begin by seeing that given some set of services J such that $u(x_i, h_{j_1}) = u(x_i, h_{j_2}) = C_0$, if we hold $\sum_{j \in J} A_{i,j} = C_1$ constant and $A_{i,k}$ constant for all $k \notin J$ such that $\sum_{k \notin J} A_{i,k} = C_2$ and $\sum_{k \notin J} A_{i,k} u(x_i, h_k) = C_3$, then any distribution of usages

between services in J doesn't affect the value of the strategic objective.

$$\begin{aligned}
L_i^u(A, H) &= \sum_{j=1}^m A_{i,j} u(x_i, h_j) - \frac{1}{q} \left(\sum_{j=1}^m A_{i,j} \right)^q \\
&= \sum_{j \in J} A_{i,j} u(x_i, h_j) + \sum_{k \notin J} A_{i,k} u(x_i, h_k) - \frac{1}{q} \left(\sum_{j \in J} A_{i,j} + \sum_{k \notin J} A_{i,k} \right)^q \\
&= C_0 \sum_{j \in J} A_{i,j} + C_3 - \frac{1}{q} \left(\sum_{j \in J} A_{i,j} + C_2 \right)^q \\
&= C_0 C_1 + C_3 - \frac{1}{q} (C_1 + C_2)^q
\end{aligned}$$

Therefore, without loss of generality, let us say all usage in J is concentrated into a service $j \in J$; since $A_{i,j'} = 0$ for all $j' \in J, j' \neq j$ we can drop them from the strategic objective and consider only a service j and services $k \notin J, k \neq j$.

Taking the derivative of $L_i^u(A, H)$ with respect to some $A_{i,j}$, we get gradient $\frac{d}{dA_{i,j}} L_i^u(A, H) = u(x_i, h_j) - \left(A_{i,j} + \sum_{k \neq j} A_{i,k} \right)^{q-1}$. Let us take two services j and k , such that $u(x_i, h_j) > u(x_i, h_l)$ for all $l \neq j$. Service k is chosen arbitrarily such that $j \neq k$. We shall now enumerate options in a case analysis focusing on the total usage of user i . We shall show that as total usage increases, the model is incentivized to concentrate all usage on the service that provides the highest utility as the derivative of the objective function with respect to other services becomes negative.

Let us say that $\sum_{l=1}^m A_{i,l} = 0$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) > 0$; therefore $A_{i,j}$ is below the optimum.

Let us say $\sum_{l=1}^m A_{i,l} < u(x_i, h_k)^{\frac{1}{q-1}}$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) - (\sum_{l=1}^m A_{i,l})^{q-1} > u(x_i, h_j) - u(x_i, h_k) > 0$; therefore $A_{i,j}$ is below the optimum. For service k , we have that $\frac{d}{dA_{i,k}} L_i^u = u(x_i, h_k) - (\sum_{l=1}^m A_{i,l})^{q-1} > u(x_i, h_k) - u(x_i, h_k) = 0$; therefore $A_{i,k}$ is below the optimum.

Let us say $\sum_{l=1}^m A_{i,l} = u(x_i, h_k)^{\frac{1}{q-1}}$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) - (\sum_{l=1}^m A_{i,l})^{q-1} = u(x_i, h_j) - u(x_i, h_k) > 0$; therefore $A_{i,j}$ is below the optimum. For service k , we have that $\frac{d}{dA_{i,k}} L_i^u = u(x_i, h_k) - (\sum_{l=1}^m A_{i,l})^{q-1} = u(x_i, h_k) - u(x_i, h_k) = 0$; therefore $A_{i,k}$ is at an optimum.

Let us say $u(x_i, h_k)^{\frac{1}{q-1}} < \sum_{l=1}^m A_{i,l} < u(x_i, h_j)^{\frac{1}{q-1}}$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) - (\sum_{l=1}^m A_{i,l})^{q-1} > u(x_i, h_j) - u(x_i, h_j) = 0$; therefore $A_{i,j}$ is below the optimum. For service k , we have that $\frac{d}{dA_{i,k}} L_i^u = u(x_i, h_k) - (\sum_{l=1}^m A_{i,l})^{q-1} < u(x_i, h_k) - u(x_i, h_k) = 0$; therefore $A_{i,k}$ is above the optimum.

Let us say $\sum_{l=1}^m A_{i,l} = u(x_i, h_j)^{\frac{1}{q-1}}$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) - (\sum_{l=1}^m A_{i,l})^{q-1} = u(x_i, h_j) - u(x_i, h_j) = 0$; therefore $A_{i,j}$ is at an optimum. For service k , we have that $\frac{d}{dA_{i,k}} L_i^u = u(x_i, h_k) - (\sum_{l=1}^m A_{i,l})^{q-1} = u(x_i, h_k) - u(x_i, h_j) < 0$; therefore $A_{i,k}$ is above the optimum.

Let us say $\sum_{l=1}^m A_{i,l} > u(x_i, h_j)^{\frac{1}{q-1}}$. For service j , we have that $\frac{d}{dA_{i,j}} L_i^u = u(x_i, h_j) - (\sum_{l=1}^m A_{i,l})^{q-1} < u(x_i, h_j) - u(x_i, h_j) = 0$; therefore $A_{i,j}$ is above the optimum. For service k , we have that $\frac{d}{dA_{i,k}} L_i^u = u(x_i, h_k) - (\sum_{l=1}^m A_{i,l})^{q-1} < u(x_i, h_k) - u(x_i, h_j) < 0$; therefore $A_{i,k}$ is above the optimum.

From all of this, we can see that the stable equilibrium holds that $A_{i,j} = u(x_i, h_j)^{\frac{1}{q-1}}$, and that for all services $k \neq j$, $A_{i,k} = 0$ because of the non-negativity of A . \square

Corollary 10. *For any user i and service j pair such that at timestep t , $u(x_i, h_j^t) \leq 0$, then $A_{i,j}^t = 0$.*

Proof of Corollary 10. By Lemma 9, we have that if there exists some service k such that $u(x_i, h_k^t) > u(x_i, h_j^t)$ then $A_{i,j}^t = 0$. Let us analyze the case where for all users k , $u(x_i, h_j^t) = 0 \geq u(x_i, h_k^t)$. Once again taking the derivative of the strategic objective for a user i (2.1) with respect to $A_{i,j}^t$ and denoting it as $\frac{d}{dA_{i,j}^t} L_i^u(A^t, H^t)$ for convenience, we can

see that if $\sum_{k=1}^m A_{i,k}^t > 0$ then $\frac{d}{dA_{i,j}^t} L_i^u(A^t, H^t) < 0$. This indicates that the optimum value of $A_{i,j}^t$ is 0. \square

Finally, we prove the main invariance result.

Proposition 3. *If a state (H^t, A^t) is zero-loss at time t , then all future states are zero-loss for all times $\tau \geq t$.*

Proof of Proposition 3. We prove the proposition by showing that if a state (H^t, A^t) is a zero-loss equilibrium, then (H^{t+1}, A^{t+1}) is also a zero-loss equilibrium.

We begin by arguing that $H^{t+1} = H^t$. Consider the retraining objective for service j at t : $\sum_{i=1}^n \frac{M_{i,j}^t}{\sum_{k=1}^n M_{k,j}^t} \ell(h_j^t, x_i, y_i)$. By Lemma 2, we know that $\ell(h_j^t, x_i, y_i) = 0$ for all users i and services j such that $M_{i,j}^{t-1} > 0$. By zero-loss condition 1 we have that $\ell(h_j^t, x_i, y_i) = 0$ for all users i and services j such that $A_{i,j}^t > 0$. Thus by the definition of memory (2.4), it must be that $\ell(h_j^t, x_i, y_i) = 0$ for all users i and services j such that $M_{i,j}^t > 0$. We conclude that h_j^t achieves zero retraining loss, and therefore by the definition of sticky tie-breaking (1), $h_j^{t+1} = h_j^t$ for all services j . This implies that the zero-loss condition 2 holds: $0 \geq u(x_i, h_j^t) = u(x_i, h_j^{t+1})$.

We next argue that if A^t is zero loss, so is A^{t+1} . First, consider negative users i with $y_i = -1$. By the zero-loss condition 2 shown in the previous paragraph $u(x_i, h_j^t) \leq 0$. Therefore, by Corollary 10 the best response is $A_{ij}^{t+1} = 0$ for all j and all negative users i , thus ensuring that zero-loss condition 1 holds for negative users i .

Now, consider positive users i with $y_i = +1$. Because $H^{t+1} = H^t$ the user best-response objective remains the same. By monotonicity of ℓ , we have that there exists some value v' such that if $\ell(h, x, +1) = 0$, $u(x, h) = v'$. As we know that $\ell(h_j^t, x_i, y_i) = 0$ for all $A_{i,j}^t > 0$, we have that for all $A_{i,j}^t > 0$, $u(x_i, h_j^t) = v'_{ij}$. Let us say that there exists some positive user i and service j such that $\ell(h_j^{t+1}, x_i, y_i) > 0$, meaning $u(x_i, h_j^{t+1}) < v'_{ij}$.

By Lemma 9, this implies that $A_{i,j}^{t+1} = 0$. As such, for all i, j , if $\ell(h_j^{t+1}, x_i, y_i) > 0$ then $A_{i,j}^{t+1} = 0$; therefore, $A^{t+1}\ell(h_j^{t+1}, x_i, y_i) = 0$ for all users i and services j . This satisfies zero-loss condition 1 for positive users.

Thus we have shown that if a state (H^t, A^t) is a zero-loss equilibrium, then (H^{t+1}, A^{t+1}) is also a zero-loss equilibrium. Through induction, this guarantees that the state (H^τ, A^τ) is a zero-loss equilibrium for all timesteps $\tau > t$. \square

A.3 Convergence

With invariance results in hand, we can now prove the main convergence result.

Lemma 5. *For any timestep t if there exists no values $M_{i,j}^{t-1} = 0$ such that $A_{i,j}^t > 0$, then (H^t, A^t) is zero-loss.*

Proof of Lemma 5. Lets say that there exists some timestep t such that there exists no values $M_{i,j}^{t-1} = 0$ such that $A_{i,j}^t > 0$.

By Lemma 2 and by the non-negativity of ℓ , this means that $\ell(h_j^t, x_i, y_i) = 0$ for all users i and services j such that $M_{i,j}^{t-1} > 0$. This gives us that for all users i and services j , $\frac{M_{i,j}^{t-1}}{\sum_{k=1}^n M_{k,j}^{t-1}} \ell(h_j^t, x_i, y_i) = 0$ as either $M_{i,j}^{t-1} = 0$ or $\ell(h_j^t, x_i, y_i) = 0$. Since $A_{i,j}^t = 0$ when $M_{i,j}^{t-1} = 0$ and when $M_{i,j}^{t-1} > 0$ it holds that $\ell(h_j^t, x_i, y_i) = 0$, we can similarly conclude that for all i, j , $A_{i,j}^t \ell(h_j^t, x_i, y_i) = 0$. This satisfies the first condition of zero-equilibrium states.

Let us define an L_i^u as follows:

$$L_i^u(A, H) = \frac{1}{q} \left(\sum_{j=1}^m A_{i,j} \right)^q - \sum_{j=1}^m A_{i,j} u(x_i, h_j)$$

We know that for all $M_{i,j}^{t-1} > 0$, $\ell(h_j^t, x_i, y_i) = 0$ by Lemma 2 and the non-negativity of ℓ . This indicates that for all i such that $y_i = -1$, if $M_{i,j}^{t-1} > 0$ then $u(x_i, h_j^t) \leq 0$. By Corollary 10, this gives that $A_{i,j}^t = 0$. Since $A_{i,j}^t = 0$ if $M_{i,j}^{t-1} = 0$ as well, this indicates that for all services j , for all users i such that $y_i = -1$, $A_{i,j}^t = 0$.

Let us say for contradiction that for some user i such that $y_i = -1$, for service $j = \operatorname{argmax}_j u(x_i, h_j^t)$ it holds that $u(x_i, h_j^t) > 0$. By Lemma 9, it must hold that $A_{i,j}^t > 0$. This poses a contradiction, and therefore we may state that for all i such that $y_i = -1$, for all j , $u(x_i, h_j^t) \leq 0$. \square

Theorem 6. *Given nonzero memory $p > 0$, there is a finite time $t \in \mathbb{N}$ after which for all $\tau > t$, (H^τ, A^τ) is zero-loss.*

Proof of Theorem 6. We may analyze each timestep t as either a timestep in which there exists some user-service pair i, j such that $A_{i,j}^t > 0$ and $M_{i,j}^{t-1} = 0$, or in which no such pair exists. There are only $n \times m$ such i, j pairs and as such a maximum of nm timesteps where the first condition is satisfied. In the contrary, at any step t where the condition isn't satisfied, by Lemma 5, we have shown that we have reached a fixed point. Therefore, there is a maximum of nm timesteps before the conditions are reached to reach a zero-loss equilibrium. By Proposition 3, this indicates that for all timesteps $\tau \geq nm$, (H^τ, A^τ) constitutes a zero-loss equilibrium. \square

We construct examples to indicate the linear dependence of convergence time on n and m . Let us define the model class as $h(x) = \operatorname{sign}(x + \theta)$, with utility as $u(x, h) = \min\{x + \theta, 1\}$ and loss as $\ell(h, x, y) = \max\{0, 1 - y(x + \theta)\}$. We set $q = 2$ and $p > 0$. User tie-breaking is done by stochastically selecting one service to use, and services tie-break by choosing the minimum change classifier between timesteps.

Example 5. Let there be n evenly spaced positive users, with features:

$$\{(0), (-0.7), (-1.4), \dots, (-0.7)(2 - n), (-0.7)(1 - n)\} : .$$

If $m = 1$ service is instantiated with model $\theta^0 = 0.5$, at every timestep t , user t will receive positive utility and elect for positive usage, pushing the classifier to $\theta^t = 1 - x_t$. This will result in $nm = n$ total timesteps before convergence.

Example 6. Let there be $n = 1$ user such that $X = \{(0)\}$ and $Y = \{-1\}$. If m services are instantiated with models $\theta_j^0 = j + 1$, at every timestep t , some service j will receive usage from the user, pushing the classifier to $\theta_j^t = -1$, at which point it'll never receive positive usage again. This will result in $nm = m$ total timesteps before convergence.

A.4 Round Robin Updates

In order to formalize the concept of round robin updates, we use the concept of rounds to encapsulate periods of time when meaningful progress towards convergence are made. With this concept, the following proof shows a finite number of such rounds will guarantee a zero-loss point.

Proposition 7. Given nonzero memory $p > 0$, there there are a finite number of rounds (Definition 4) $r \in \mathbb{N}$ after which for all $\rho > r$, (H^ρ, A^ρ) is zero-loss.

Proof of Proposition 7. We shall prove this by showing that a round is functionally equivalent to a set of timesteps that assume stochastic user tie-breaking. Since we have that a zero-loss point will be reached after a finite number of timesteps by Theorem 6, this will imply the existence of a zero-loss point after a finite number of rounds.

As we have already analyzed the result of a service best responding to a set of users,

and a user best responding to a set of services, let us analyze the change when services and users undergo multiple consecutive updates.

Given a service j at update $k > 0$, let us denote h_j^k as the best response to memory M_j^k updating on usages A . By the memory update (2.4), we have that for all users i , $M_{i,j}^k = 0$ if and only if $M_{i,j}^{k+1} = 0$. By Lemma 2, we have that $M_{i,j}^{k+1} \ell(h_j^k, x_i, y_i) = 0$ for all users i . As h_j^k achieves a value of zero for the objective value of the service update, by the non-negativity of loss and M we have that h_j^k is a best-response to M^{k+1} . Sticky updating gives that $h_j^{k+1} = h_j^k$.

Given a user i at update $k > 0$, let us denote A_i^k as the best response to services H . A_i^{k+1} will be a user best response to H ; since no other variables are involved in the user objective function, this is equivalent to choosing a new sample A_i^k from the equivalence class of usages that maximizes the user objective function on H .

Due to the separability of the joint user and service updates, these asynchronous updates can be reanalyzed as parts of the joint update. As such, we can collapse all consecutive user updates and all consecutive service updates; without loss of generality, rounds can now be seen as a series of alternating user and service updates. This can be re-indexed as a series of timesteps, each composed of one joint user and one joint service update.

By Theorem 6, this concludes the proof. \square

A.5 Non-Realizability

Proposition 8. *When there doesn't exist any $h \in H$ such that $\ell(h, x_i, y_i) = 0$ for all $i = 1, \dots, n$, there exist initializations in which the state (H, A) never reaches a stable*

state.

Proof of Proposition 8. The memoryless case is trivial, as we may repeat the proof of Proposition 1 and simply add a positive point $x = (-1, -1)$, $y = -1$ that would never be used by the services but nevertheless violates non-realizability.

Let there exist a set of users with features $X = \{0.3, -0.3\}$ and labels $Y = \{1, -1\}$, interacting in a game with one service, where we set memory parameter $p = 1$. The service model class, user utility, and service loss are all provided below. When the service provides exactly zero utility, users tie-break by not using the service. The service updates with memory

$$\mathcal{H} = \left\{ h(x) = \begin{cases} +1 & x > \theta \\ -1 & x \leq \theta \end{cases} \text{ s.t. } \theta \in \mathbb{R} \right\}$$

$$u(x, h) = \mathbf{1}\{x - \theta > 0\}$$

$$\ell(h, x, y) = \max\{1 - y(x - \theta), 0\} + \mathbf{1}\{y(x - \theta) \leq 0\}$$

Suppose that the initial service model is defined through parameter $\theta = 0$. Through analysis of optimal updates, we can see that the below updates would occur.

$\theta^0 = 0$	$A^0 = \begin{bmatrix} 1 & 0 \end{bmatrix}^\top$	$M^0 = \begin{bmatrix} 0.5 & 0 \end{bmatrix}^\top$
$\theta^1 = -0.7$	$A^1 = \begin{bmatrix} 1 & 1 \end{bmatrix}^\top$	$M^1 = \begin{bmatrix} 0.75 & 0.5 \end{bmatrix}^\top$
$\theta^2 = -0.3$	$A^2 = \begin{bmatrix} 1 & 0 \end{bmatrix}^\top$	$M^2 = \begin{bmatrix} 0.875 & 0.25 \end{bmatrix}^\top$
$\theta^3 = -0.3$	$A^3 = \begin{bmatrix} 1 & 0 \end{bmatrix}^\top$	$M^3 = \begin{bmatrix} 0.9375 & 0.125 \end{bmatrix}^\top$
$\theta^4 = -0.7$	$A^4 = \begin{bmatrix} 1 & 1 \end{bmatrix}^\top$	$M^4 = \begin{bmatrix} 0.96875 & 0.5625 \end{bmatrix}^\top$
$\theta^5 = -0.3$	$A^5 = \begin{bmatrix} 1 & 1 \end{bmatrix}^\top$	$M^5 = \begin{bmatrix} 0.984375 & 0.28125 \end{bmatrix}^\top$

This will oscillate endlessly between $\theta = -0.3$ and $\theta = -0.7$, never resulting in a stable state as $M_{0,1}$ will always decrease to where the negative point is given utility. \square

References

- [1] Guy Aridor, Yishay Mansour, Aleksandrs Slivkins, and Zhiwei Steven Wu. Competing bandits: The perils of exploration under competition. *arXiv preprint arXiv:2007.10144*, 2020.
- [2] Omer Ben-Porat and Moshe Tennenholtz. Best response regression. *Advances in Neural Information Processing Systems*, 30, 2017.
- [3] Omer Ben-Porat and Moshe Tennenholtz. Regression equilibrium. In *Proceedings of the 2019 ACM Conference on Economics and Computation*, pages 173–191, 2019.
- [4] Yatong Chen, Jialu Wang, and Yang Liu. Strategic recourse in linear classification. *ArXiv*, abs/2011.00355, 2020.
- [5] Jennifer Chien, Margaret Roberts, and Berk Ustun. Algorithmic censoring in dynamic learning systems. *arXiv preprint arXiv:2305.09035*, 2023.
- [6] Sarah Dean, Mihaela Curmei, Lillian J Ratliff, Jamie Morgenstern, and Maryam Fazel. Multi-learner risk reduction under endogenous participation dynamics. *arXiv preprint arXiv:2206.02667*, 2022.
- [7] Jinshuo Dong, Aaron Roth, Zachary Schutzman, Bo Waggoner, and Zhiwei Steven Wu. Strategic classification from revealed preferences. In *Proceedings of the 2018 ACM Conference on Economics and Computation*, pages 55–70, 2018.
- [8] Eilin Francis, Joshua Blumenstock, and Jonathan Robinson. Digital credit: A snapshot of the current landscape and open research questions. *CEGA White Paper*, pages 1739–76, 2017.

- [9] Ganesh Ghalme, Vineet Nair, Itay Eilat, Inbal Talgam-Cohen, and Nir Rosenfeld. Strategic classification in the dark. In *International Conference on Machine Learning*, pages 3672–3681. PMLR, 2021.
- [10] Tony Ginart, Eva Zhang, Yongchan Kwon, and James Zou. Competing ai: How does competition feedback affect machine learning? In *International Conference on Artificial Intelligence and Statistics*, pages 1693–1701. PMLR, 2021.
- [11] Ronen Gradwohl and Moshe Tennenholtz. Coopetition against an amazon. In *International Symposium on Algorithmic Game Theory*, pages 347–365. Springer, 2022.
- [12] Moritz Hardt, Meena Jagadeesan, and Celestine Mendler-Dünnér. Performative power. *arXiv preprint arXiv:2203.17232*, 2022.
- [13] Moritz Hardt, Nimrod Megiddo, Christos Papadimitriou, and Mary Wootters. Strategic classification. In *Proceedings of the 2016 ACM conference on innovations in theoretical computer science*, pages 111–122, 2016.
- [14] Keegan Harris, Chara Podimata, and Zhiwei Steven Wu. Strategic apple tasting. *arXiv preprint arXiv:2306.06250*, 2023.
- [15] Tatsunori Hashimoto, Megha Srivastava, Hongseok Namkoong, and Percy Liang. Fairness without demographics in repeated loss minimization. In *International Conference on Machine Learning*, pages 1929–1938. PMLR, 2018.
- [16] Lily Hu, Nicole Immorlica, and Jennifer Wortman Vaughan. The disparate effects of strategic manipulation. In *Proceedings of the Conference on Fairness, Accountability, and Transparency, FAT* ’19*, page 259–268, New York, NY, USA, 2019. Association for Computing Machinery.

- [17] Meena Jagadeesan, Michael I Jordan, and Nika Haghtalab. Competition, alignment, and equilibria in digital marketplaces. *arXiv preprint arXiv:2208.14423*, 2022.
- [18] Sérgio Jesus, José Pombal, Duarte Alves, André Cruz, Pedro Saleiro, Rita P. Ribeiro, João Gama, and Pedro Bizarro. Turning the tables: Biased, imbalanced, dynamic tabular datasets for ml evaluation, 2022.
- [19] Jon Kleinberg and Manish Raghavan. How do classifiers induce agents to invest effort strategically? *ACM Transactions on Economics and Computation (TEAC)*, 8(4):1–23, 2020.
- [20] Yongchan Kwon, Antonio Ginart, and James Zou. Competition over data: how does data purchase affect users? *arXiv preprint arXiv:2201.10774*, 2022.
- [21] Volker Lohweg. banknote authentication. UCI Machine Learning Repository, 2013. DOI: <https://doi.org/10.24432/C55P57>.
- [22] John Miller, Smitha Milli, and Moritz Hardt. Strategic classification is causal modeling in disguise. In *International Conference on Machine Learning*, pages 6917–6926. PMLR, 2020.
- [23] Smitha Milli, John Miller, Anca D Dragan, and Moritz Hardt. The social cost of strategic classification. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 230–239, 2019.
- [24] Vineet Nair, Ganesh Ghalme, Inbal Talgam-Cohen, and Nir Rosenfeld. Strategic representation, 2022.
- [25] Adhyayan Narang, Evan Faulkner, Dmitriy Drusvyatskiy, Maryam Fazel, and Lillian J Ratliff. Multiplayer performative prediction: Learning in decision-dependent games. *arXiv preprint arXiv:2201.03398*, 2022.

- [26] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [27] Juan Perdomo, Tijana Zrnic, Celestine Mendler-Dünnér, and Moritz Hardt. Performative prediction. In *International Conference on Machine Learning*, pages 7599–7609. PMLR, 2020.
- [28] Georgios Piliouras and Fang-Yi Yu. Multi-agent performative prediction: From global stability and optimality to chaos. *arXiv preprint arXiv:2201.10483*, 2022.
- [29] Killian Wood and Emiliano Dall’Anese. Stochastic saddle point problems with decision-dependent distributions. *arXiv preprint arXiv:2201.02313*, 2022.
- [30] Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding deep learning requires rethinking generalization, 2017.
- [31] Xueru Zhang, Mohammadmahdi Khaliligarekani, Cem Tekin, et al. Group retention when using machine learning in sequential decision making: the interplay between user dynamics and fairness. *Advances in Neural Information Processing Systems*, 32, 2019.
- [32] Tijana Zrnic, Eric Mazumdar, Shankar Sastry, and Michael Jordan. Who leads and who follows in strategic classification? *Advances in Neural Information Processing Systems*, 34:15257–15269, 2021.